

How to monitor your micro-service with Prometheus?



Wojciech Barczyński - SMACC.io | [LI](#) | [Twitter](#) | 16 October 2018

ABOUT ME

- Lead Software Developer- SMACC (FinTech/AI)
- Before:
System Engineer && Developer Lyke
- Looking:
Tools for efficient and :) teams

POINT OF VIEW

- Software Developer
- startups && fast-moving environment
- Infra and platform should just work
- Infra and platform ~ invisible

MY GOAL

- Start with monitoring
- Metrics? Start with Up/Down -> USE || RED
- Drop-in solution - Prometheus
- How to start? With this demo app

**WHY?
MONOLIT ;)**



WHY? MICROSERVICES ;)



OBSERVABILITY

- Monitoring
- Logging
- Tracing

OBSERVABILITY

	Metrics	Logging	Tracing
CapEx	Medium	Low	High
OpEx	Low	High	Medium
Reaction	High	Medium	Low
Investigation	Low	Medium	High

Go for Industrial Programming by Peter Bourgon

NOT A SILVER-BULLET

but:

- Easy to setup
- Immediately value

Surprisingly: the last one implemented

CENTRALIZED LOGGING

- Like a debugging vs testing
- Usually much too late
- Post-mortem
- Hard to find the needle

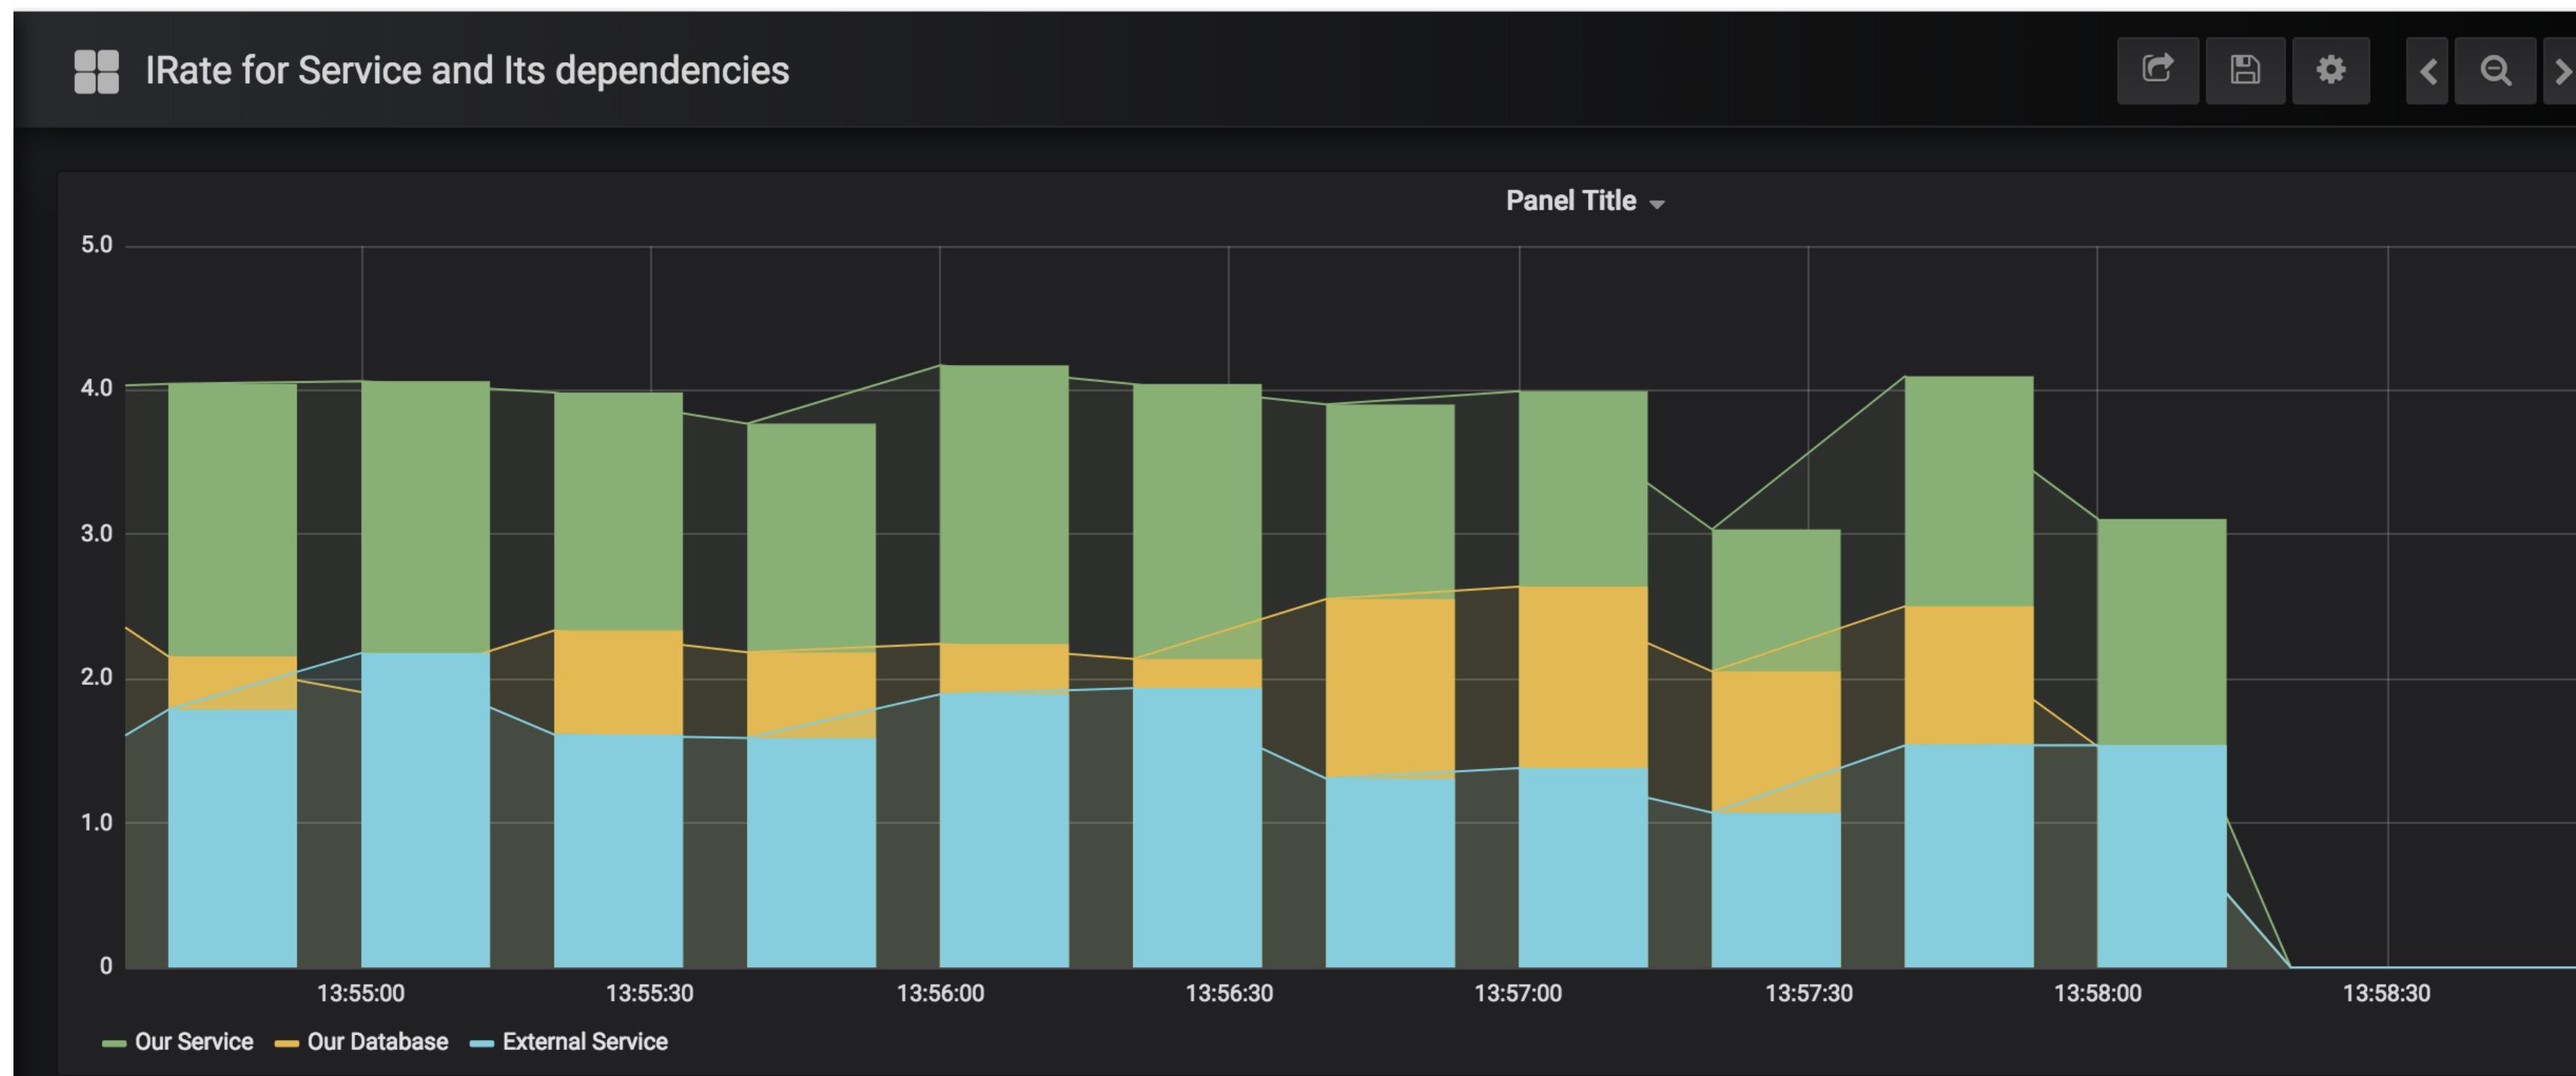
MONITORING

- Numbers
- Trends
- Dependencies
- + Actions

METRIC

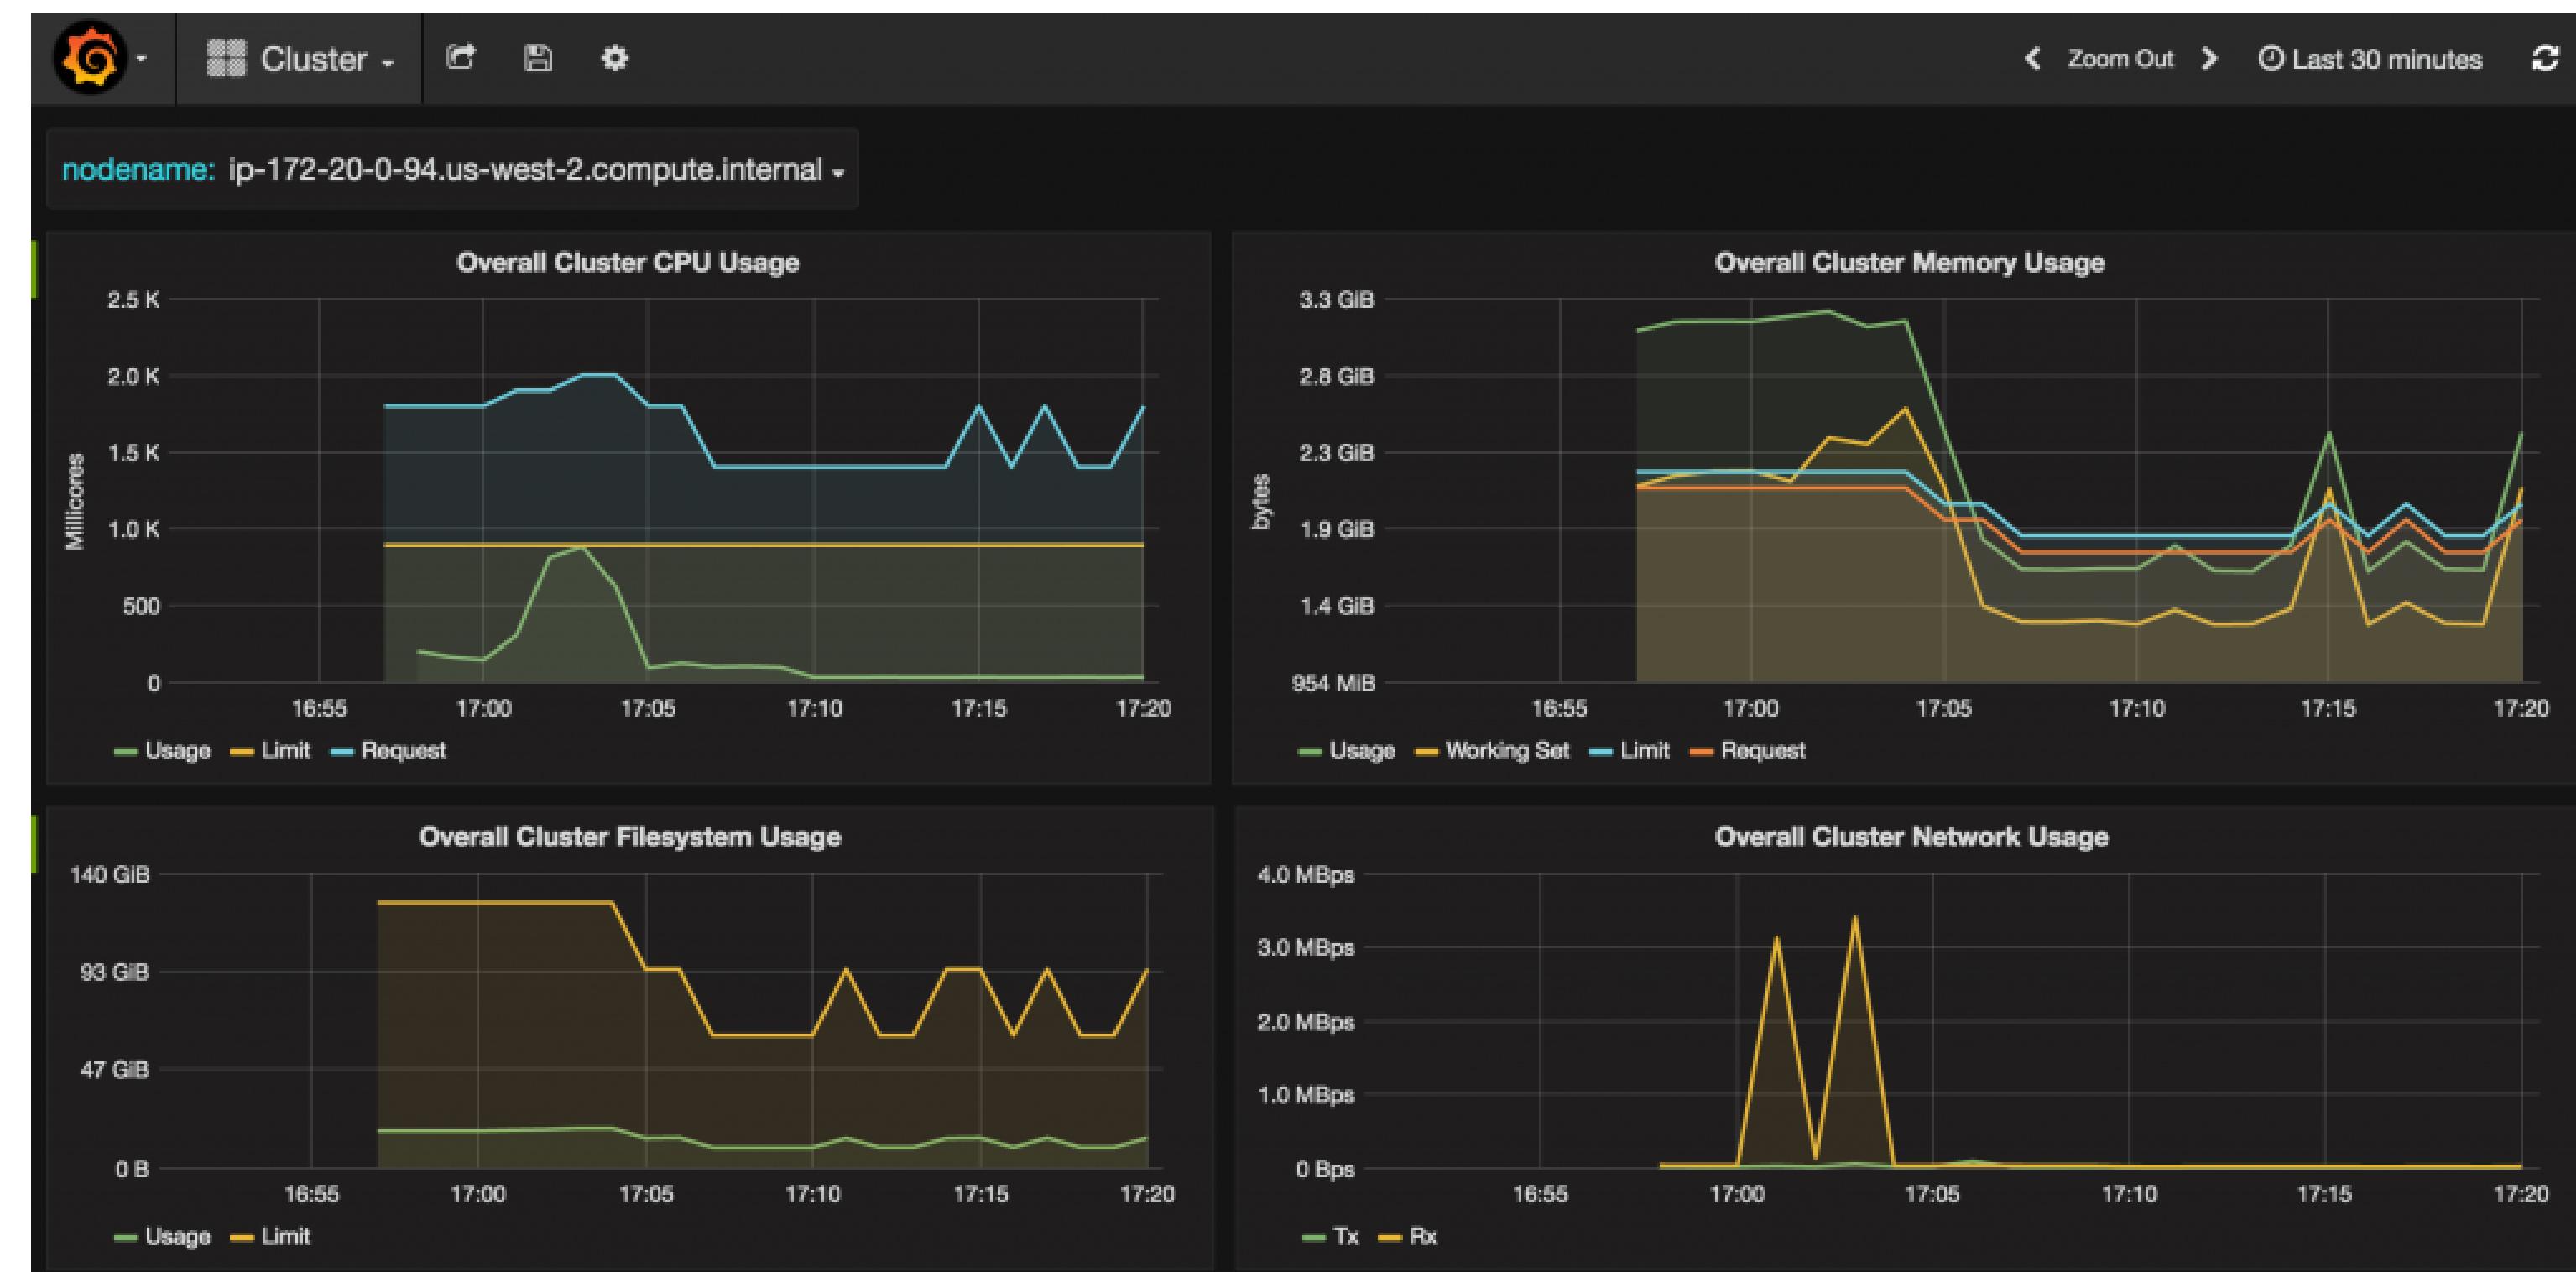
Name	Label	Value
traefik_requests_total	code="200", method="GET"	3001

MONITORING



Demo app

MONITORING



Example from [couchbase blog](#)

HOW TO FIND THE RIGHT METRIC?

HOW TO FIND THE RIGHT METRIC?

- USE
- RED

USE

Utilization the average time that the resource was busy servicing work

Saturation extra work which it can't service, often queued

Errors the count of error events

Documented and Promoted by [Berdan Gregg](#)

USE

- Utilization: as a percent over a time interval: "one disk is running at 90% utilization".
- Saturation:
- Errors:

USE

- **Utilization:**
- **Saturation:** as a queue length. eg, "the CPUs have an average run queue length of four".
- **Errors:**

USE

- **utilization:**
- **saturation:**
- **errors:** scalar counts. eg, "this network interface drops packages".

USE

- traditionally more instance oriented
- still useful in the microservices world

RED

Rate

How busy is your service?

Error

Errors

Duration

What is the latency of my service?

Tom Wilkie's guideline for instrumenting applications.

RED

- **Rate** - how many requests per seconds handled
- **Error**
- **Duration (distribution)**

RED

- **Rate**
- **Error** - how many request per seconds handled we failed
- **Duration**

RED

- Rate
- Error
- Duration - how long the requests took

RED

- Follow Four Golden Signals by Google SREs [1]
- Focus on what matters for end-users

[1] Latency, Traffic, Errors, Saturation ([src](#))

RED

Not recommended for:

- batch-oriented
- streaming services

PROMETHEUS



WHAT PROMETHEUS IS?

- Aggregation of time-series data
- Not an event-based system

PROMETHEUS STACK

- Prometheus - collect
- Alertmanager - alerts
- Grafana - visualize

PROMETHEUS

- Plain text
- Metrics collected over HTTP *metrics*/[1]
- Pull model [2]
- Wide support for languages
- PromQL
- Integration with Kubernetes, Traefik, ...

[1] most common, *prometheus* in Spring 1.5

[2] See *scrape time*, push-mode possible

METRICS IN PLAIN TEXT

```
# HELP order_mgmt_audit_duration_seconds Multiprocess metric
# TYPE order_mgmt_audit_duration_seconds summary
order_mgmt_audit_duration_seconds_count{status_code="200"} 41.
order_mgmt_audit_duration_seconds_sum{status_code="200"} 27.44
order_mgmt_audit_duration_seconds_count{status_code="500"} 1.0
order_mgmt_audit_duration_seconds_sum{status_code="500"} 0.716
# HELP order_mgmt_duration_seconds Multiprocess metric
# TYPE order_mgmt_duration_seconds summary
order_mgmt_duration_seconds_count{method="GET",path="/complex"}
order_mgmt_duration_seconds_sum{method="GET",path="/complex",s}
order_mgmt_duration_seconds_count{method="GET",path="/",status}
order_mgmt_duration_seconds_sum{method="GET",path="/",status_c}
order_mgmt_duration_seconds_count{method="GET",path="/complex"}
order_mgmt_duration_seconds_sum{method="GET",path="/complex",s}
```

METRICS IN PLAIN TEXT

```
# HELP go_gc_duration_seconds A summary of the GC invocation d
# TYPE go_gc_duration_seconds summary
go_gc_duration_seconds{quantile="0"} 9.01e-05
go_gc_duration_seconds{quantile="0.25"} 0.000141101
go_gc_duration_seconds{quantile="0.5"} 0.000178902
go_gc_duration_seconds{quantile="0.75"} 0.000226903
go_gc_duration_seconds{quantile="1"} 0.006099658
go_gc_duration_seconds_sum 18.749046756
go_gc_duration_seconds_count 89273
```

EXPORTERS

- Mongodb
- Mysql
- Postgresql
- Rabbitmq
- ...
- also Blackbox exporter

Examples: [memcached](#), [psql](#)

PROMETHEUS PromQL

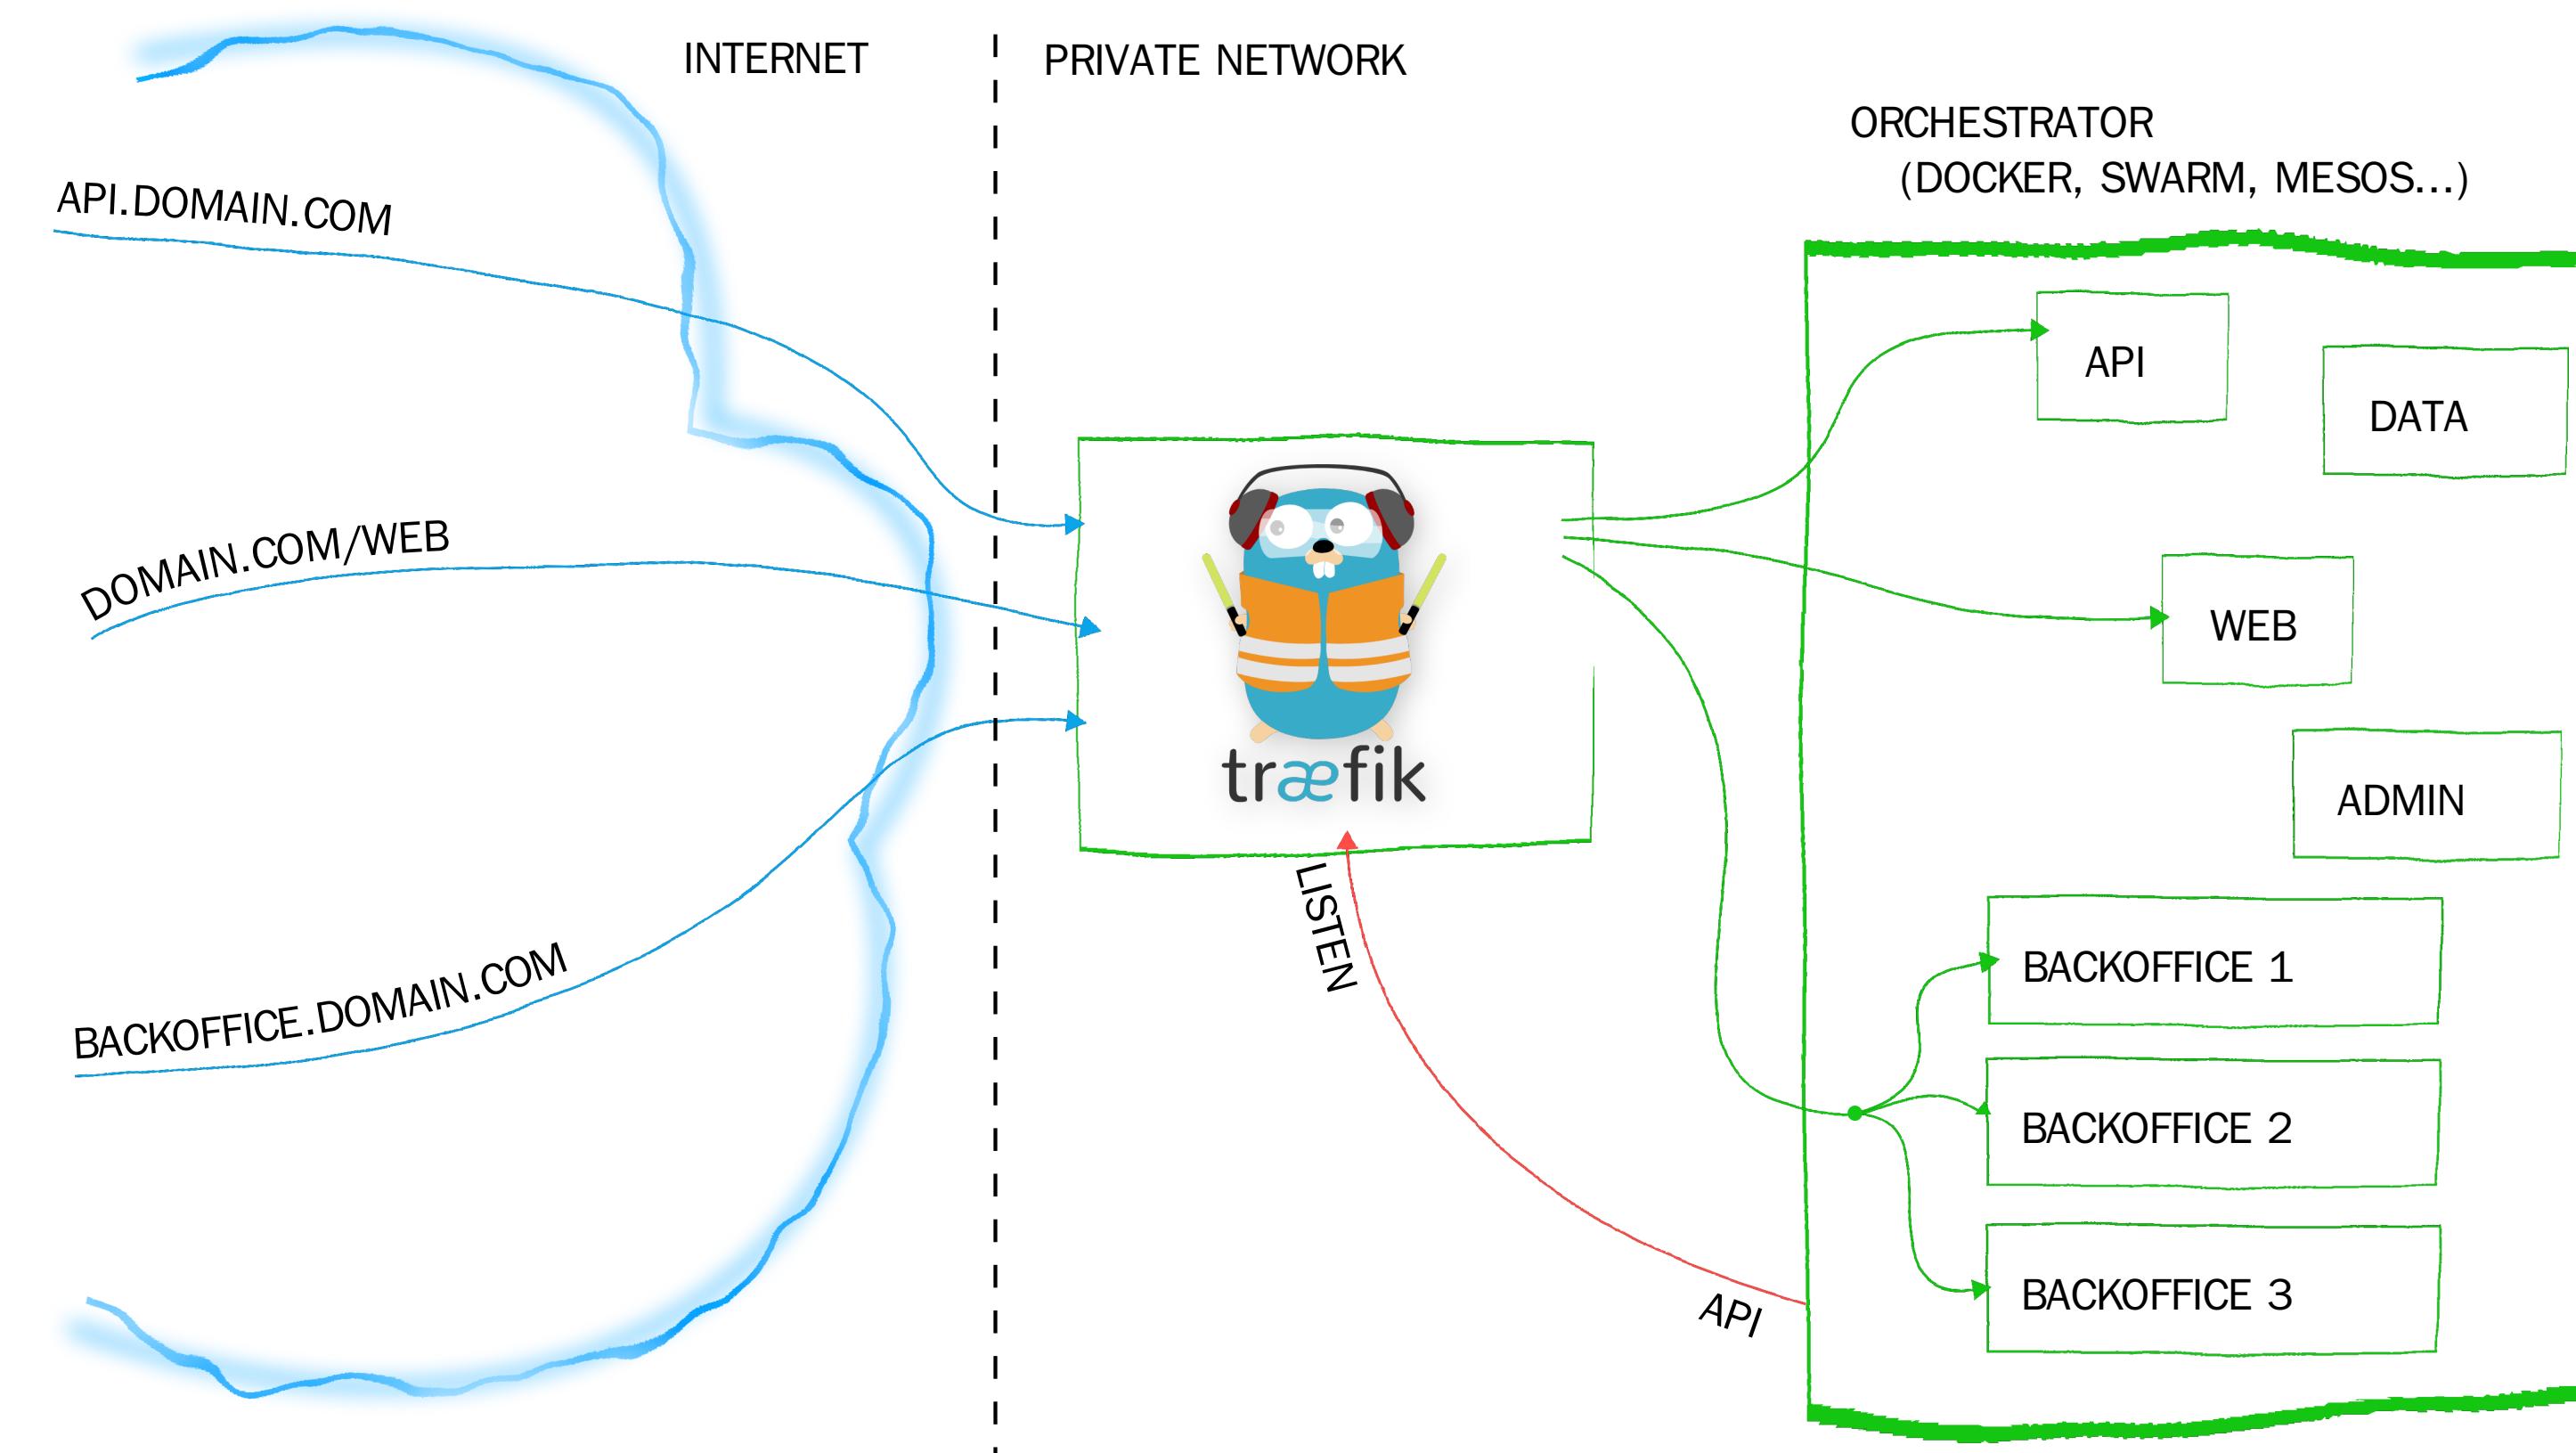
```
# working with histograms:  
  
histogram_quantile(0.9,  
    rate(http_req_duration_seconds_bucket[10m]) )  
  
# rates:  
  
rate(http_requests_total{job="api-server"} [5m] )  
irate(http_requests_total{job="api-server"} [5m] )  
  
# more complex:  
  
redict_linear()  
holt_winters()
```

PROMETHEUS PromQL

Alarming:

```
ALERT ProductionAppServiceInstanceDown
  IF up { environment = "production", app =~ ".+" } == 0
  FOR 4m
  ANNOTATIONS {
    summary = "Instance of {{$labels.app}} is down",
    description = " Instance {{$labels.instance}} of app
  }
```

CLOUD-NATIVE PROJECTS INTEGRATION



- --web.metrics.prometheus

METRICS

- Counter - just up
- Gauge - up/down
- Histogram
- Summary

HISTOGRAM

traefik_duration_seconds_bucket

{method="GET, code="200" }

{le="0.1"}	2229
{le="0.3"}	107
{le="1.2"}	100
{le="5"}	4
{le="+Inf"}	2
_sum	
_count	2342

SUMMARY

http_request_duration_seconds

{quantile="0.5"}	4
------------------	---

{quantile="0.9"}	5
------------------	---

http_request_duration_seconds_sum	9
-----------------------------------	---

http_request_duration_seconds_count	3
-------------------------------------	---

HISTOGRAM / SUMMARY:

- Latency of services
- Request or Request size

Histograms recommended

RED

Metric + PromQL:

```
sum(irate(order_mgmt_duration_seconds_count  
{job=~".*"} [1m])) by (status_code)
```

METRIC AND LABEL NAMING

Best practises on **metric names**:

- service name is your prefix `user_`
- state the base unit `_seconds` and `_bytes`

PROMETHEUS + JAVA



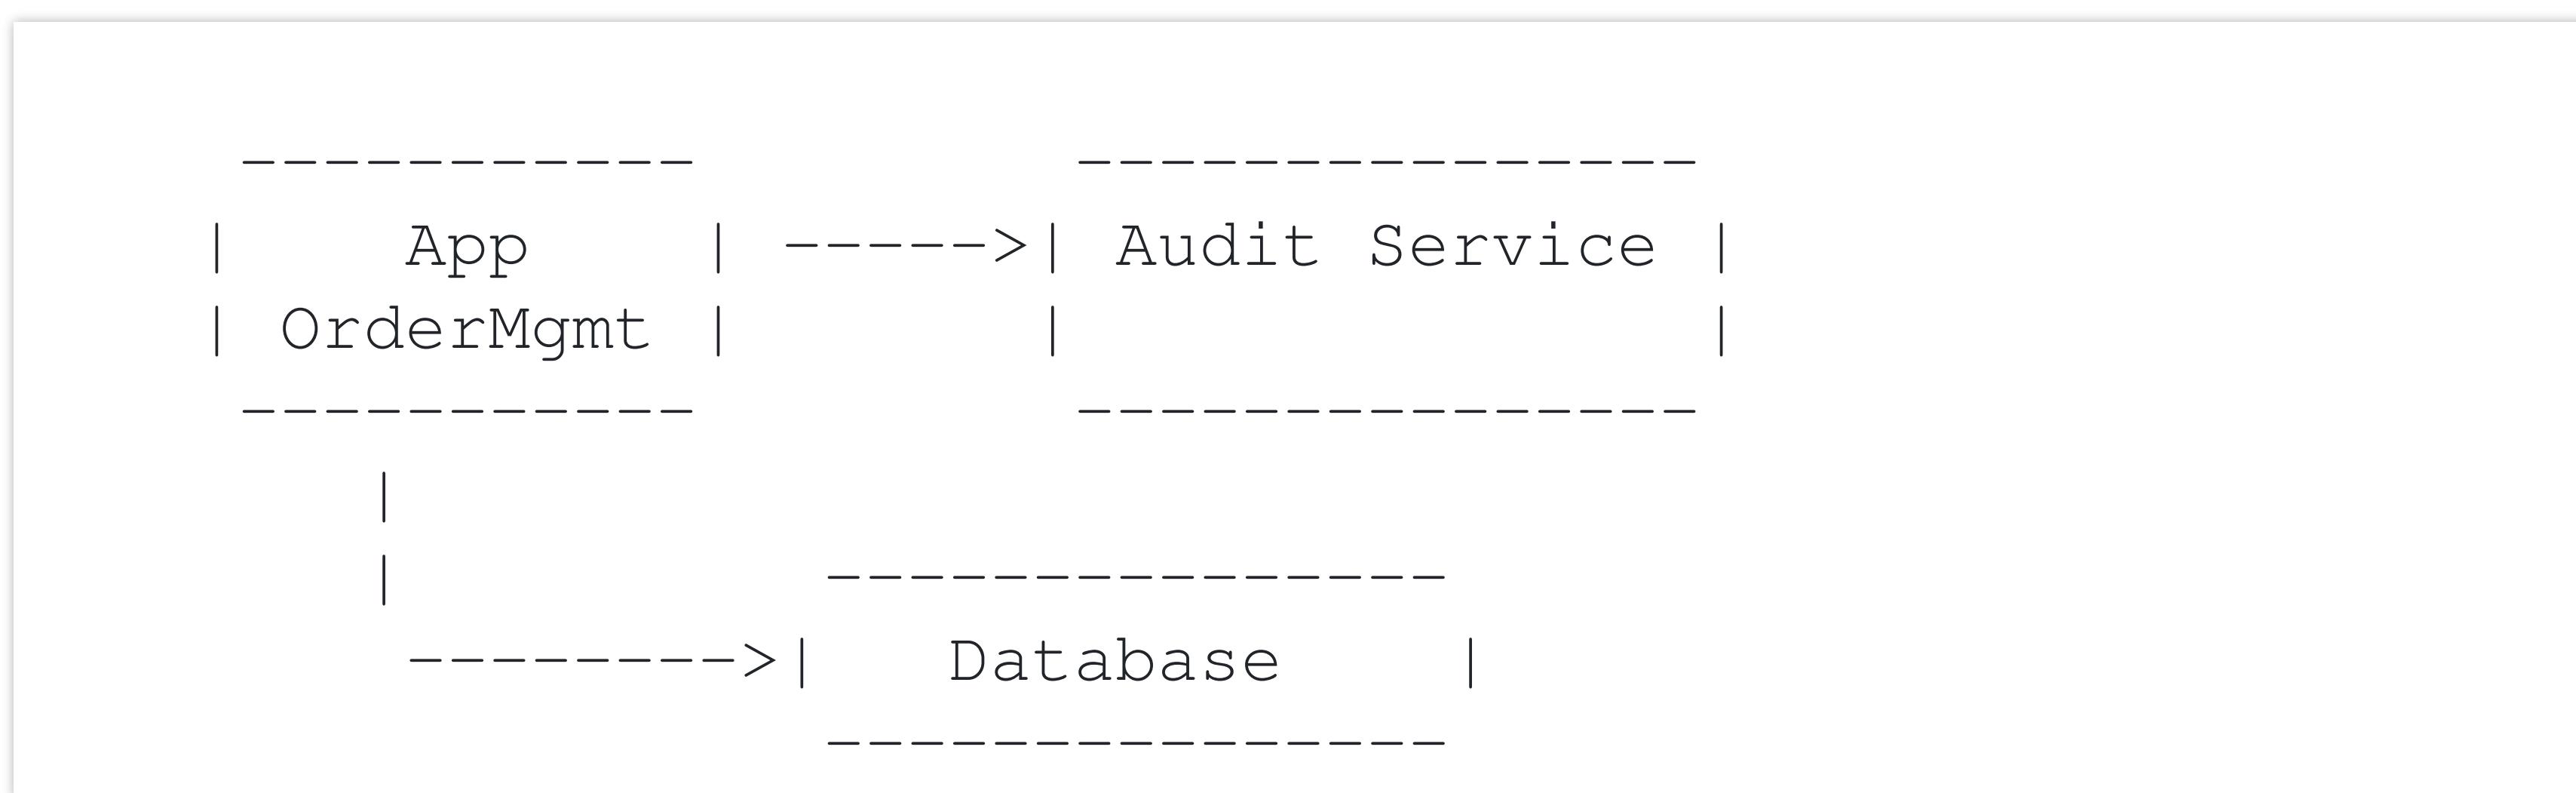
JAVA CLIENT

- client_python
- Counter
- Gauge
- Summary
- Histogram

JAVA

- Sprint 1.5 - demo
- Sprint 2.x - wrapper with `io.micrometer.core`
- JMX exporter -
https://github.com/prometheus/jmx_exporter

DEMO: SIMPLE REST SERVICE



DEMO:

- <http://127.0.0.1:8080/hello> - service
- <http://127.0.0.1:8080/prometheus/>
- <http://127.0.0.1:9090> - prometheus
- <http://127.0.0.1:3000> - grafana
- <http://127.0.0.1:9093> - alertmanager

DEMO

```
☁ demo $ make start
```

```
☁ demo $ docker ps
```

CONTAINER ID	IMAGE	PORTS
5f824d1bc789	grafana/grafana:5.2.2	0.0.0.0:3000->3
d681a414a8b6	prom/prometheus:v2.1.0	0.0.0.0:9090->9
ea0d9233e159	prom/alertmanager:v0.15.1	0.0.0.0:9093->9
732c59fb3753	java-prom_order-manager	0.0.0.0:8080->8

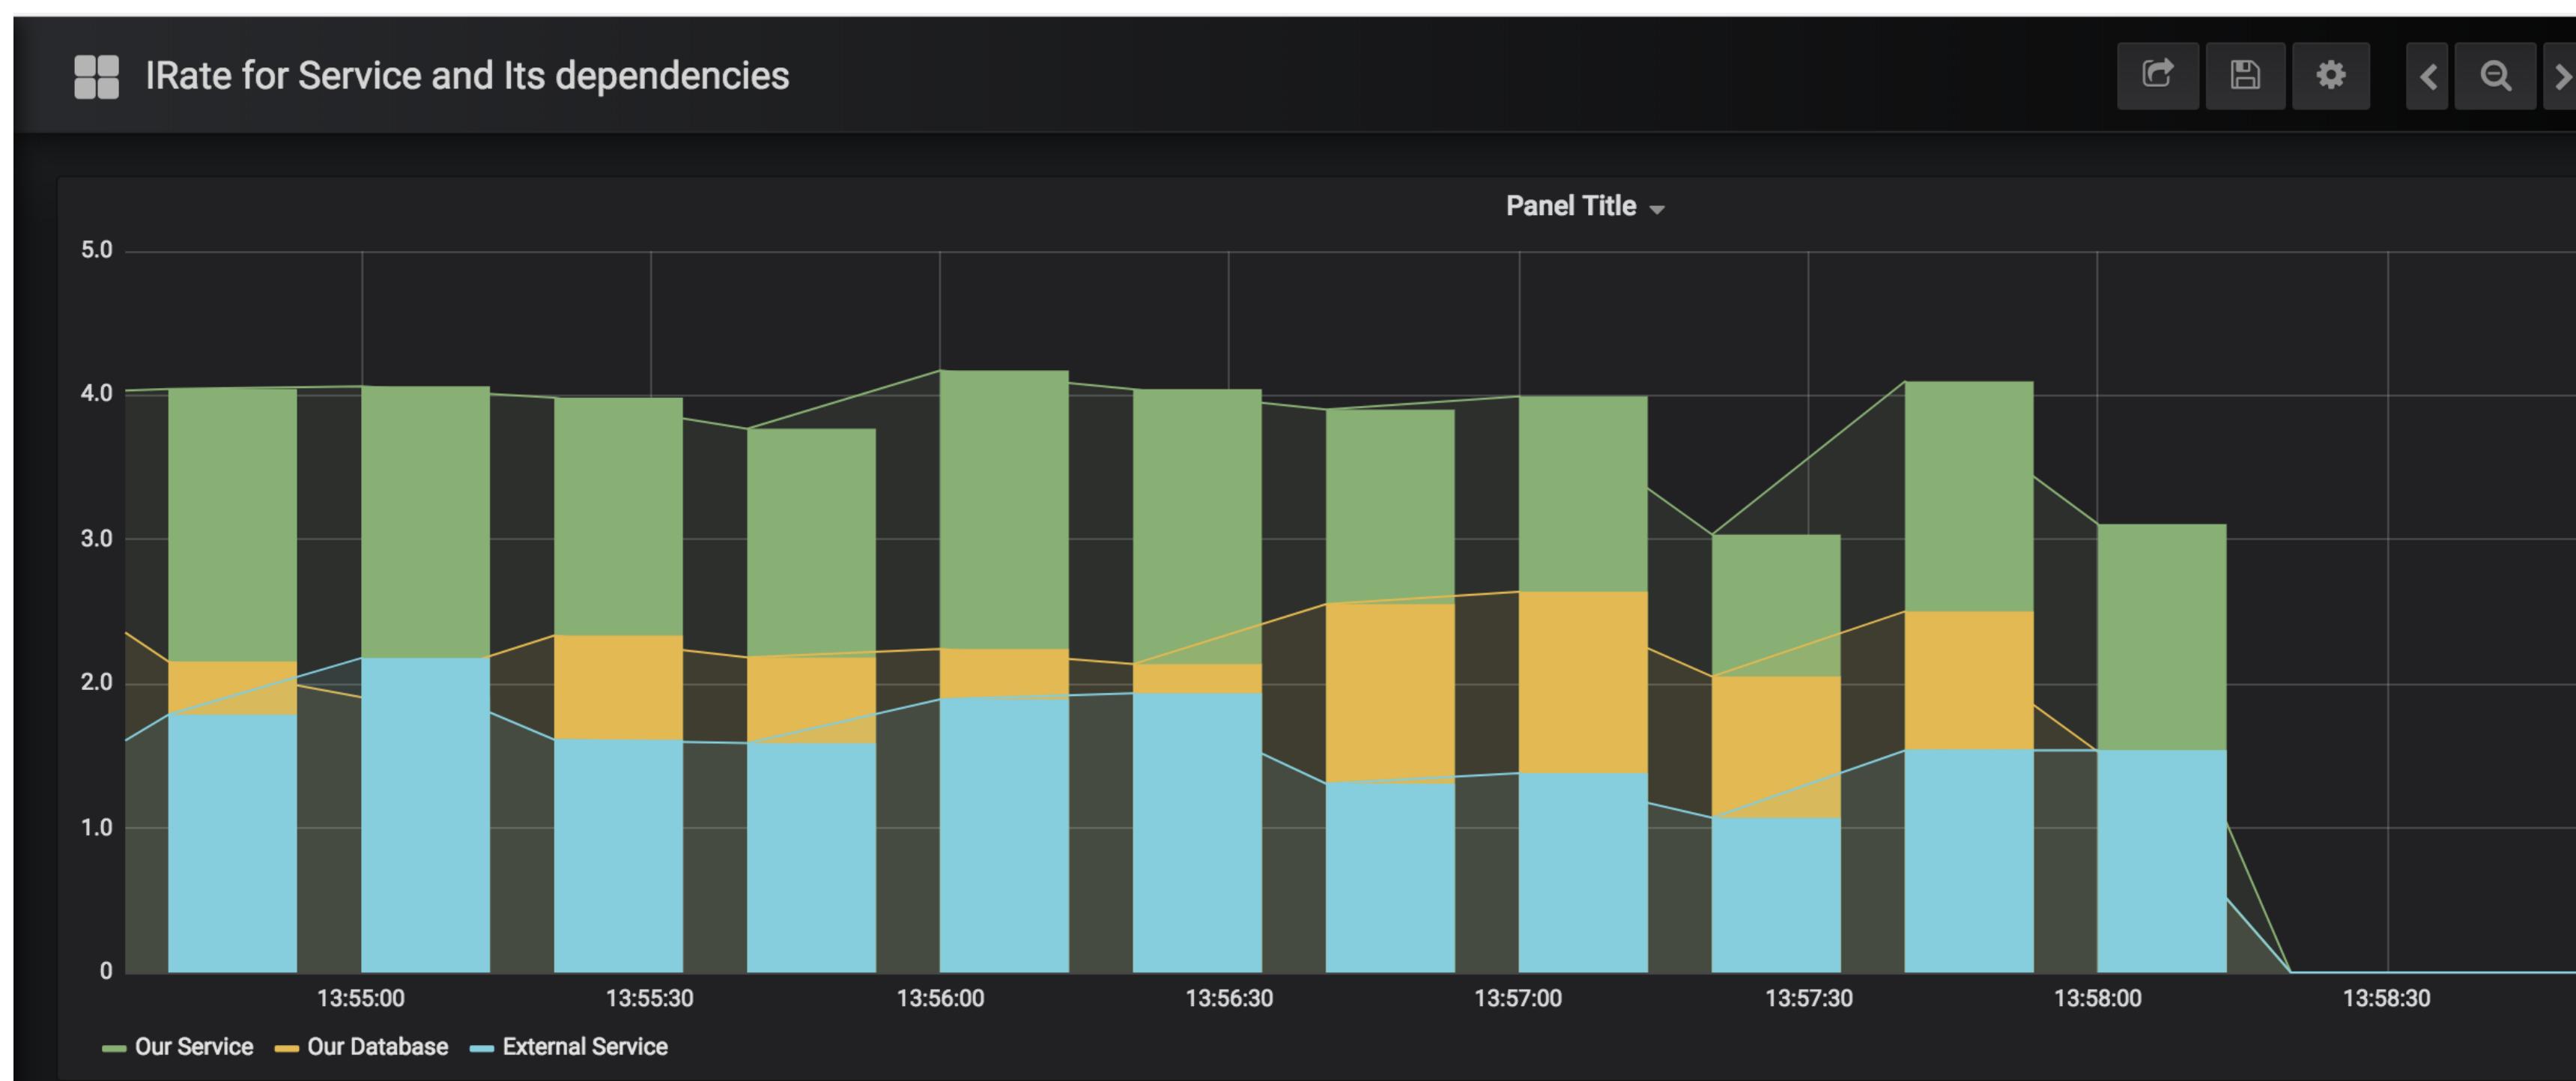
DEMO: GENERATE CALLS

```
▶ demo ⚡ make srv_wrk_random
```

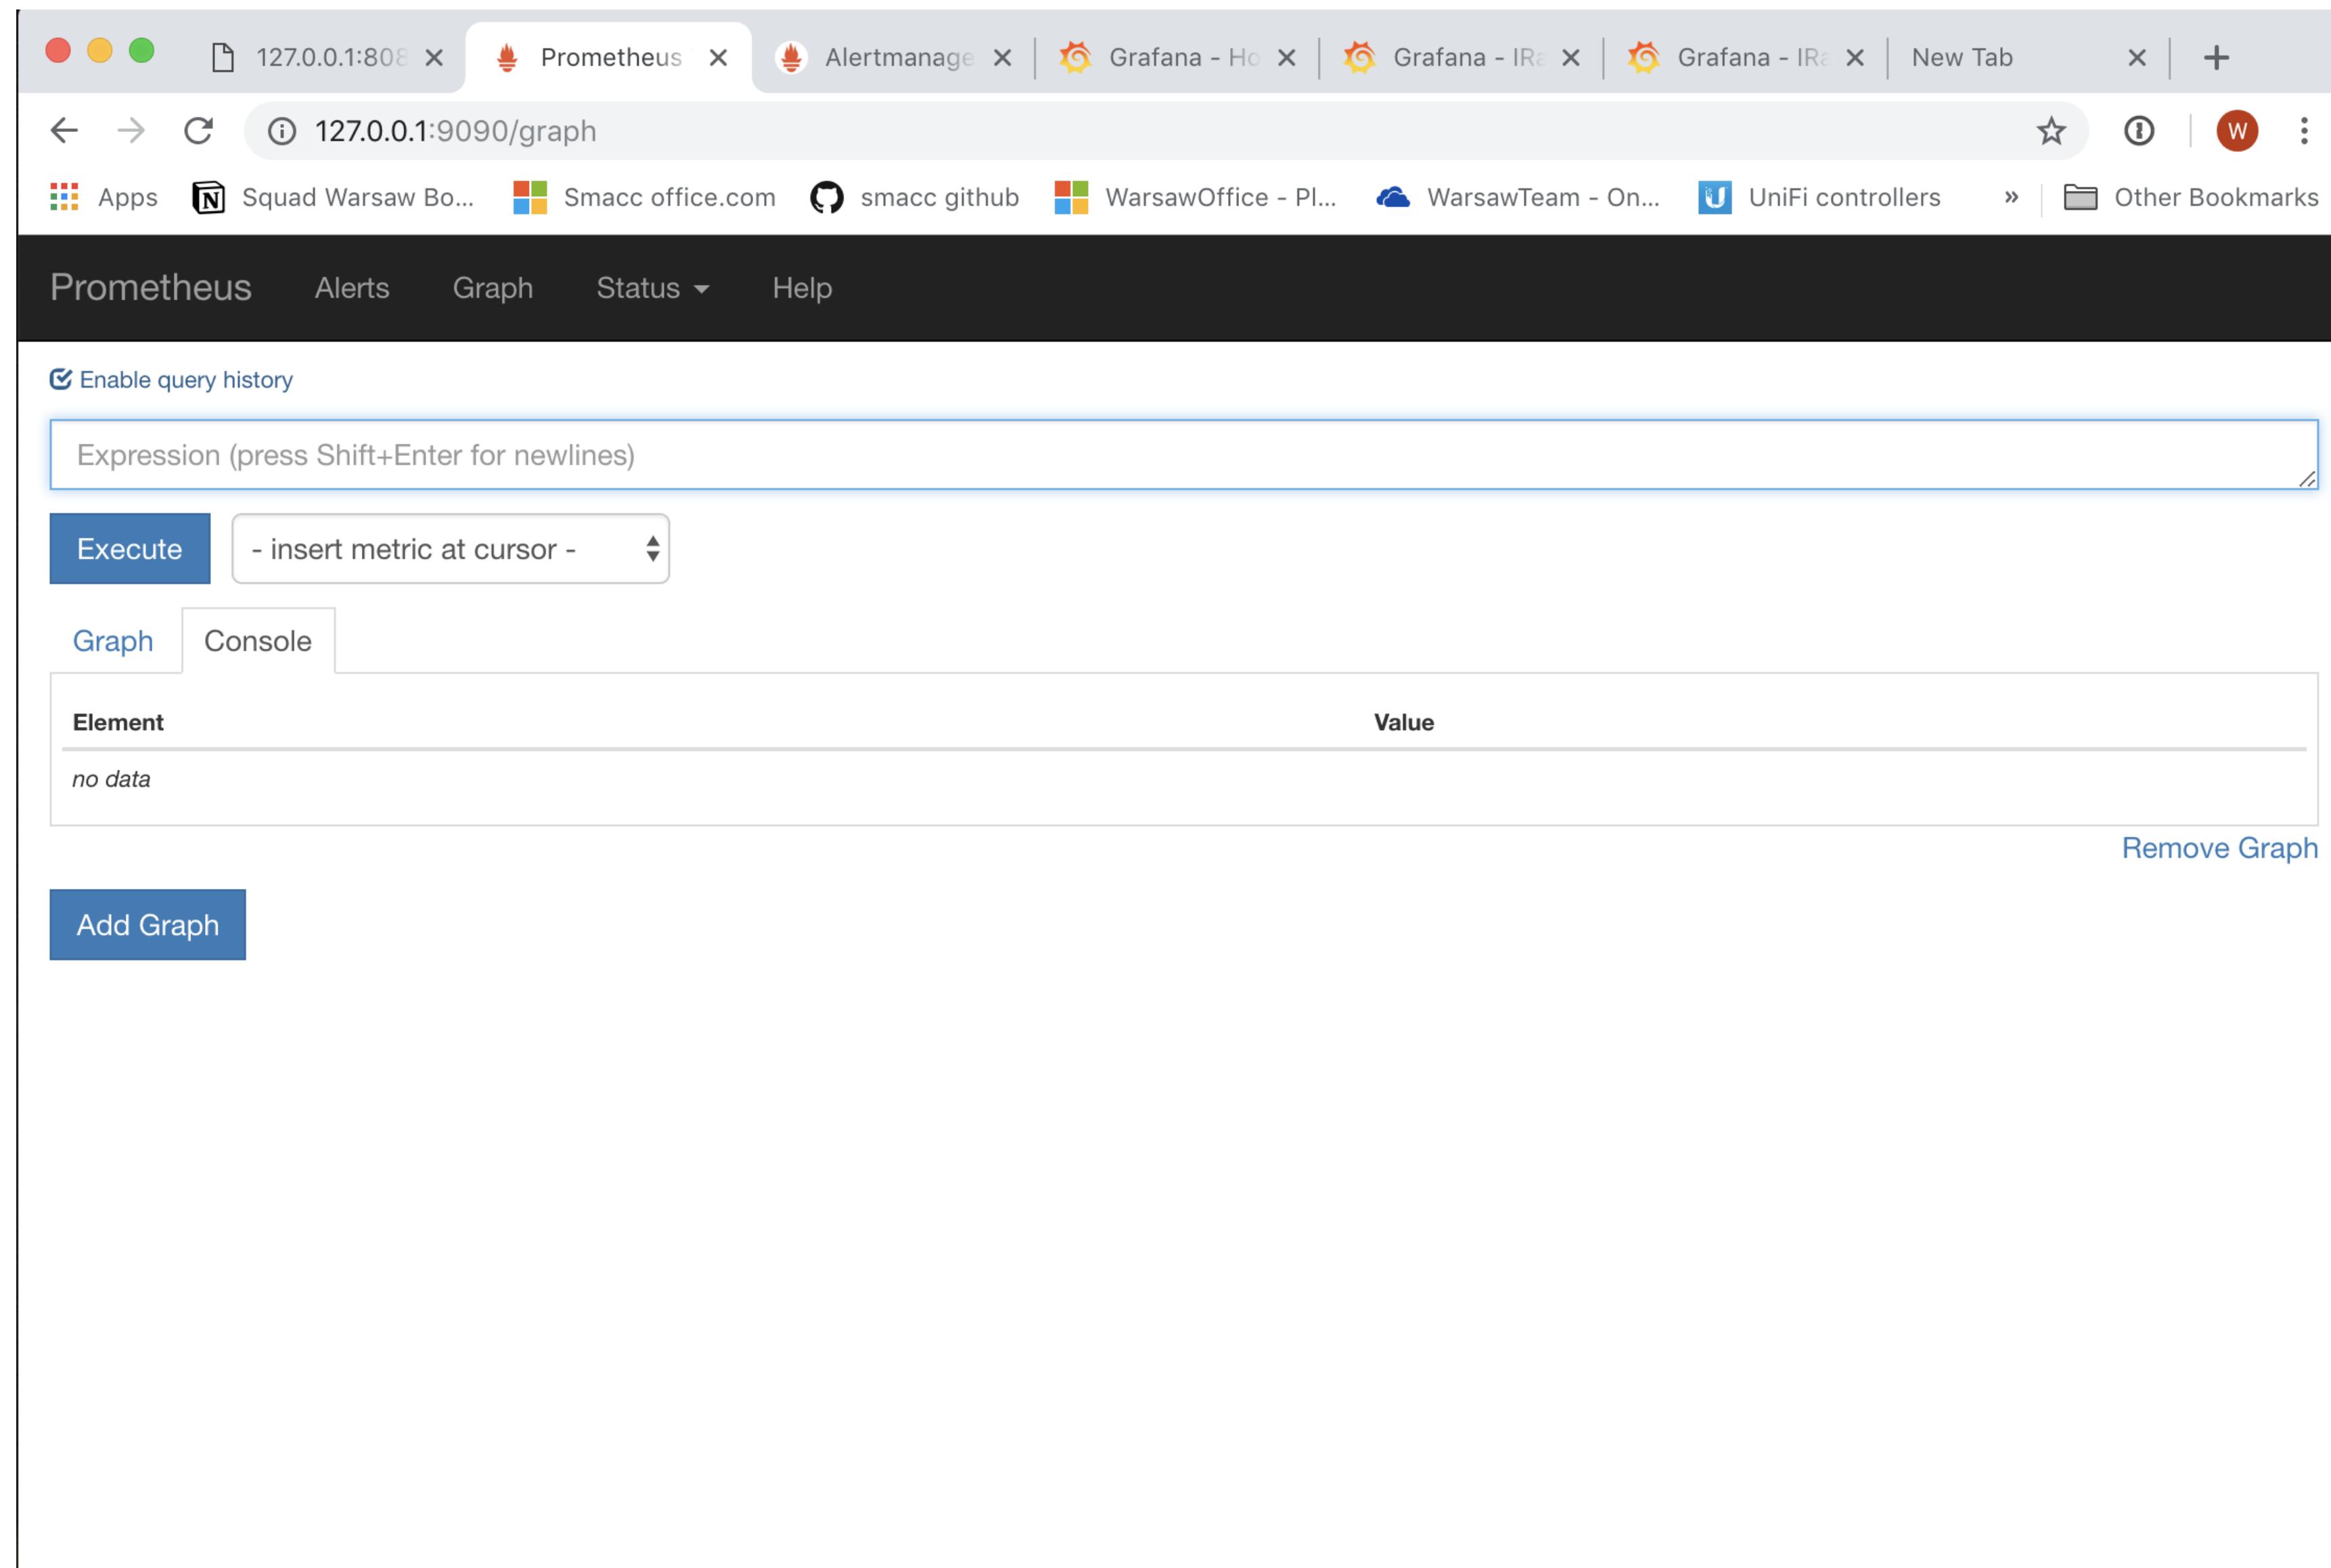
With error injection

HELP order_mgmt_database_duration_seconds Multiprocess metric
TYPE order_mgmt_database_duration_seconds histogram
order_mgmt_database_duration_seconds_sum{sql_state="0",status_code="0"} 338.9096608161926
order_mgmt_database_duration_seconds_sum{sql_state="HY000",status_code="1001"} 92.89293241500854
order_mgmt_database_duration_seconds_bucket{le="0.1",sql_state="0",status_code="0"} 72.0
order_mgmt_database_duration_seconds_bucket{le="0.25",sql_state="0",status_code="0"} 168.0
order_mgmt_database_duration_seconds_bucket{le="0.5",sql_state="0",status_code="0"} 287.0
order_mgmt_database_duration_seconds_bucket{le="0.75",sql_state="0",status_code="0"} 459.0
order_mgmt_database_duration_seconds_bucket{le="0.9",sql_state="0",status_code="0"} 573.0
order_mgmt_database_duration_seconds_bucket{le="1.0",sql_state="0",status_code="0"} 646.0
order_mgmt_database_duration_seconds_bucket{le="2.5",sql_state="0",status_code="0"} 648.0
order_mgmt_database_duration_seconds_bucket{le="+Inf",sql_state="0",status_code="0"} 648.0
order_mgmt_database_duration_seconds_count{sql_state="0",status_code="0"} 648.0
order_mgmt_database_duration_seconds_bucket{le="0.1",sql_state="HY000",status_code="1001"} 24.0
order_mgmt_database_duration_seconds_bucket{le="0.25",sql_state="HY000",status_code="1001"} 86.0
order_mgmt_database_duration_seconds_bucket{le="0.5",sql_state="HY000",status_code="1001"} 134.0
order_mgmt_database_duration_seconds_bucket{le="0.75",sql_state="HY000",status_code="1001"} 190.0
order_mgmt_database_duration_seconds_bucket{le="0.9",sql_state="HY000",status_code="1001"} 217.0
order_mgmt_database_duration_seconds_bucket{le="1.0",sql_state="HY000",status_code="1001"} 225.0
order_mgmt_database_duration_seconds_bucket{le="2.5",sql_state="HY000",status_code="1001"} 225.0
order_mgmt_database_duration_seconds_bucket{le="+Inf",sql_state="HY000",status_code="1001"} 225.0
order_mgmt_database_duration_seconds_count{sql_state="HY000",status_code="1001"} 225.0
HELP order_mgmt_audit_duration_seconds Multiprocess metric
TYPE order_mgmt_audit_duration_seconds summary
order_mgmt_audit_duration_seconds_count{status_code="200"} 490.0
order_mgmt_audit_duration_seconds_sum{status_code="200"} 231.30700039863586
order_mgmt_audit_duration_seconds_count{status_code="500"} 158.0
order_mgmt_audit_duration_seconds_sum{status_code="500"} 99.42281174659729
HELP order_mgmt_duration_seconds Multiprocess metric
TYPE order_mgmt_duration_seconds summary
order_mgmt_duration_seconds_count{method="GET",path="/complex",status_code="200"} 490.0
order_mgmt_duration_seconds_sum{method="GET",path="/complex",status_code="200"} 471.54110622406006
order_mgmt_duration_seconds_count{method="GET",path="/complex",status_code="503"} 383.0
order_mgmt_duration_seconds_sum{method="GET",path="/complex",status_code="503"} 291.444673538208

GRAFANA



PROMETHEUS



PROMETHEUS

The screenshot shows the Prometheus web interface at `127.0.0.1:9090/graph`. A dropdown menu is open over a list of metrics. The menu includes options like "Execute", "Graph", "Element", and "Add Graph". The "Graph" option is currently selected. The dropdown lists several metrics, with "order_mgmt_audit_duration_seconds_count" highlighted in red.

Expression (press Shift+Enter for newlines)

Execute ✓ - insert metric at cursor -
Graph ALERTS
Element order_mgmt_audit_duration_seconds_count
no data order_mgmt_audit_duration_seconds_sum
order_mgmt_database_duration_seconds_bucket
order_mgmt_database_duration_seconds_count
order_mgmt_database_duration_seconds_sum
order_mgmt_duration_seconds_count
order_mgmt_duration_seconds_sum
scrape_duration_seconds
scrape_samples_post_metric_relabeling
scrape_samples_scraped
up

Value Remove Graph

PROMETHEUS

The screenshot shows a web browser window for the Prometheus interface at `127.0.0.1:9090/alerts`. The browser's tab bar includes tabs for Prometheus, Alertmanager, Grafana, and New Tab. The main content area is titled "Alerts".

ProductionInstanceDown (0 active)

```
alert: ProductionInstanceDown
expr: up{env="production"} == 0
for: 2m
labels:
  severity: opsgenie
annotations:
  description: '{{ $labels.instance }} of job {{ $labels.job }} has been down for more than 2 minutes.'
  summary: Instance {{ $labels.instance }} of {{ $labels.app }} is down
```

StagingAppServiceInstanceDown (0 active)

KILL THE SERVICE

```
☁ demo ✚ docker stop java-prom_order-manager_1
```

PROMETHEUS

The screenshot shows the Prometheus web interface at `127.0.0.1:9090/alerts`. The page title is "Alerts".

ProductionInstanceDown (1 active)

```
alert: ProductionInstanceDown
expr: up{env="production"} == 0
for: 2m
labels:
  severity: opsgenie
annotations:
  description: '{{ $labels.instance }} of job {{ $labels.job }} has been down for more than 2 minutes.'
  summary: Instance {{ $labels.instance }} of {{ $labels.app }} is down
```

Labels	State	Active Since	Value
alertname="ProductionInstanceDown" app="order-manager" env="production" instance="order-manager:8080" job="my-service" owner="wb@example.com" severity="opsgenie" tier="front-end"	PENDING	2018-10-02 09:35:20.5392318 +0000 UTC	0

StagingAppServiceInstanceDown (0 active)

PROMETHEUS

The screenshot shows a web browser window for the Prometheus interface at `127.0.0.1:9090/alerts`. The browser's address bar also lists other tabs for Prometheus, Alertmanager, Grafana, and New Tab.

The main content area is titled "Alerts". It displays one active alert named "ProductionInstanceDown".

ProductionInstanceDown (1 active)

```
alert: ProductionInstanceDown
expr: up{env="production"} == 0
for: 2m
labels:
  severity: opsgenie
annotations:
  description: '{{ $labels.instance }} of job {{ $labels.job }} has been down for more than 2 minutes.'
  summary: Instance {{ $labels.instance }} of {{ $labels.app }} is down
```

Labels

Labels	State	Active Since	Value
<code>alertname="ProductionInstanceDown"</code> <code>app="order-manager"</code> <code>env="production"</code> <code>instance="order-manager:8080"</code> <code>job="my-service"</code> <code>owner="wb@example.com"</code> <code>severity="opsgenie"</code> <code>tier="front-end"</code>	FIRING	2018-10-02 09:35:20.5392318 +0000 UTC	0

StagingAppServiceInstanceDown (0 active)

ALERTMANAGER

The screenshot shows the Alertmanager interface running locally at `127.0.0.1:9093/#/alerts`. The top navigation bar includes tabs for Prometheus, Alertmanager, Grafana, and New Tab. Below the header, there are links for Apps, Squad Warsaw Bo..., Smacc office.com, smacc github, WarsawOffice - Pl..., WarsawTeam - On..., UniFi controllers, and Other Bookmarks.

The main content area has tabs for Alertmanager, Alerts, Silences, and Status, with "Alertmanager" selected. A "New Silence" button is located in the top right corner of this section.

Below the tabs, there are two filter options: "Filter" and "Group". The "Receiver: All" dropdown is set to "All". There are also checkboxes for "Silenced" and "Inhibited". A "Custom matcher, e.g. `env="production"`" input field is present with a "+" button to add more matchers.

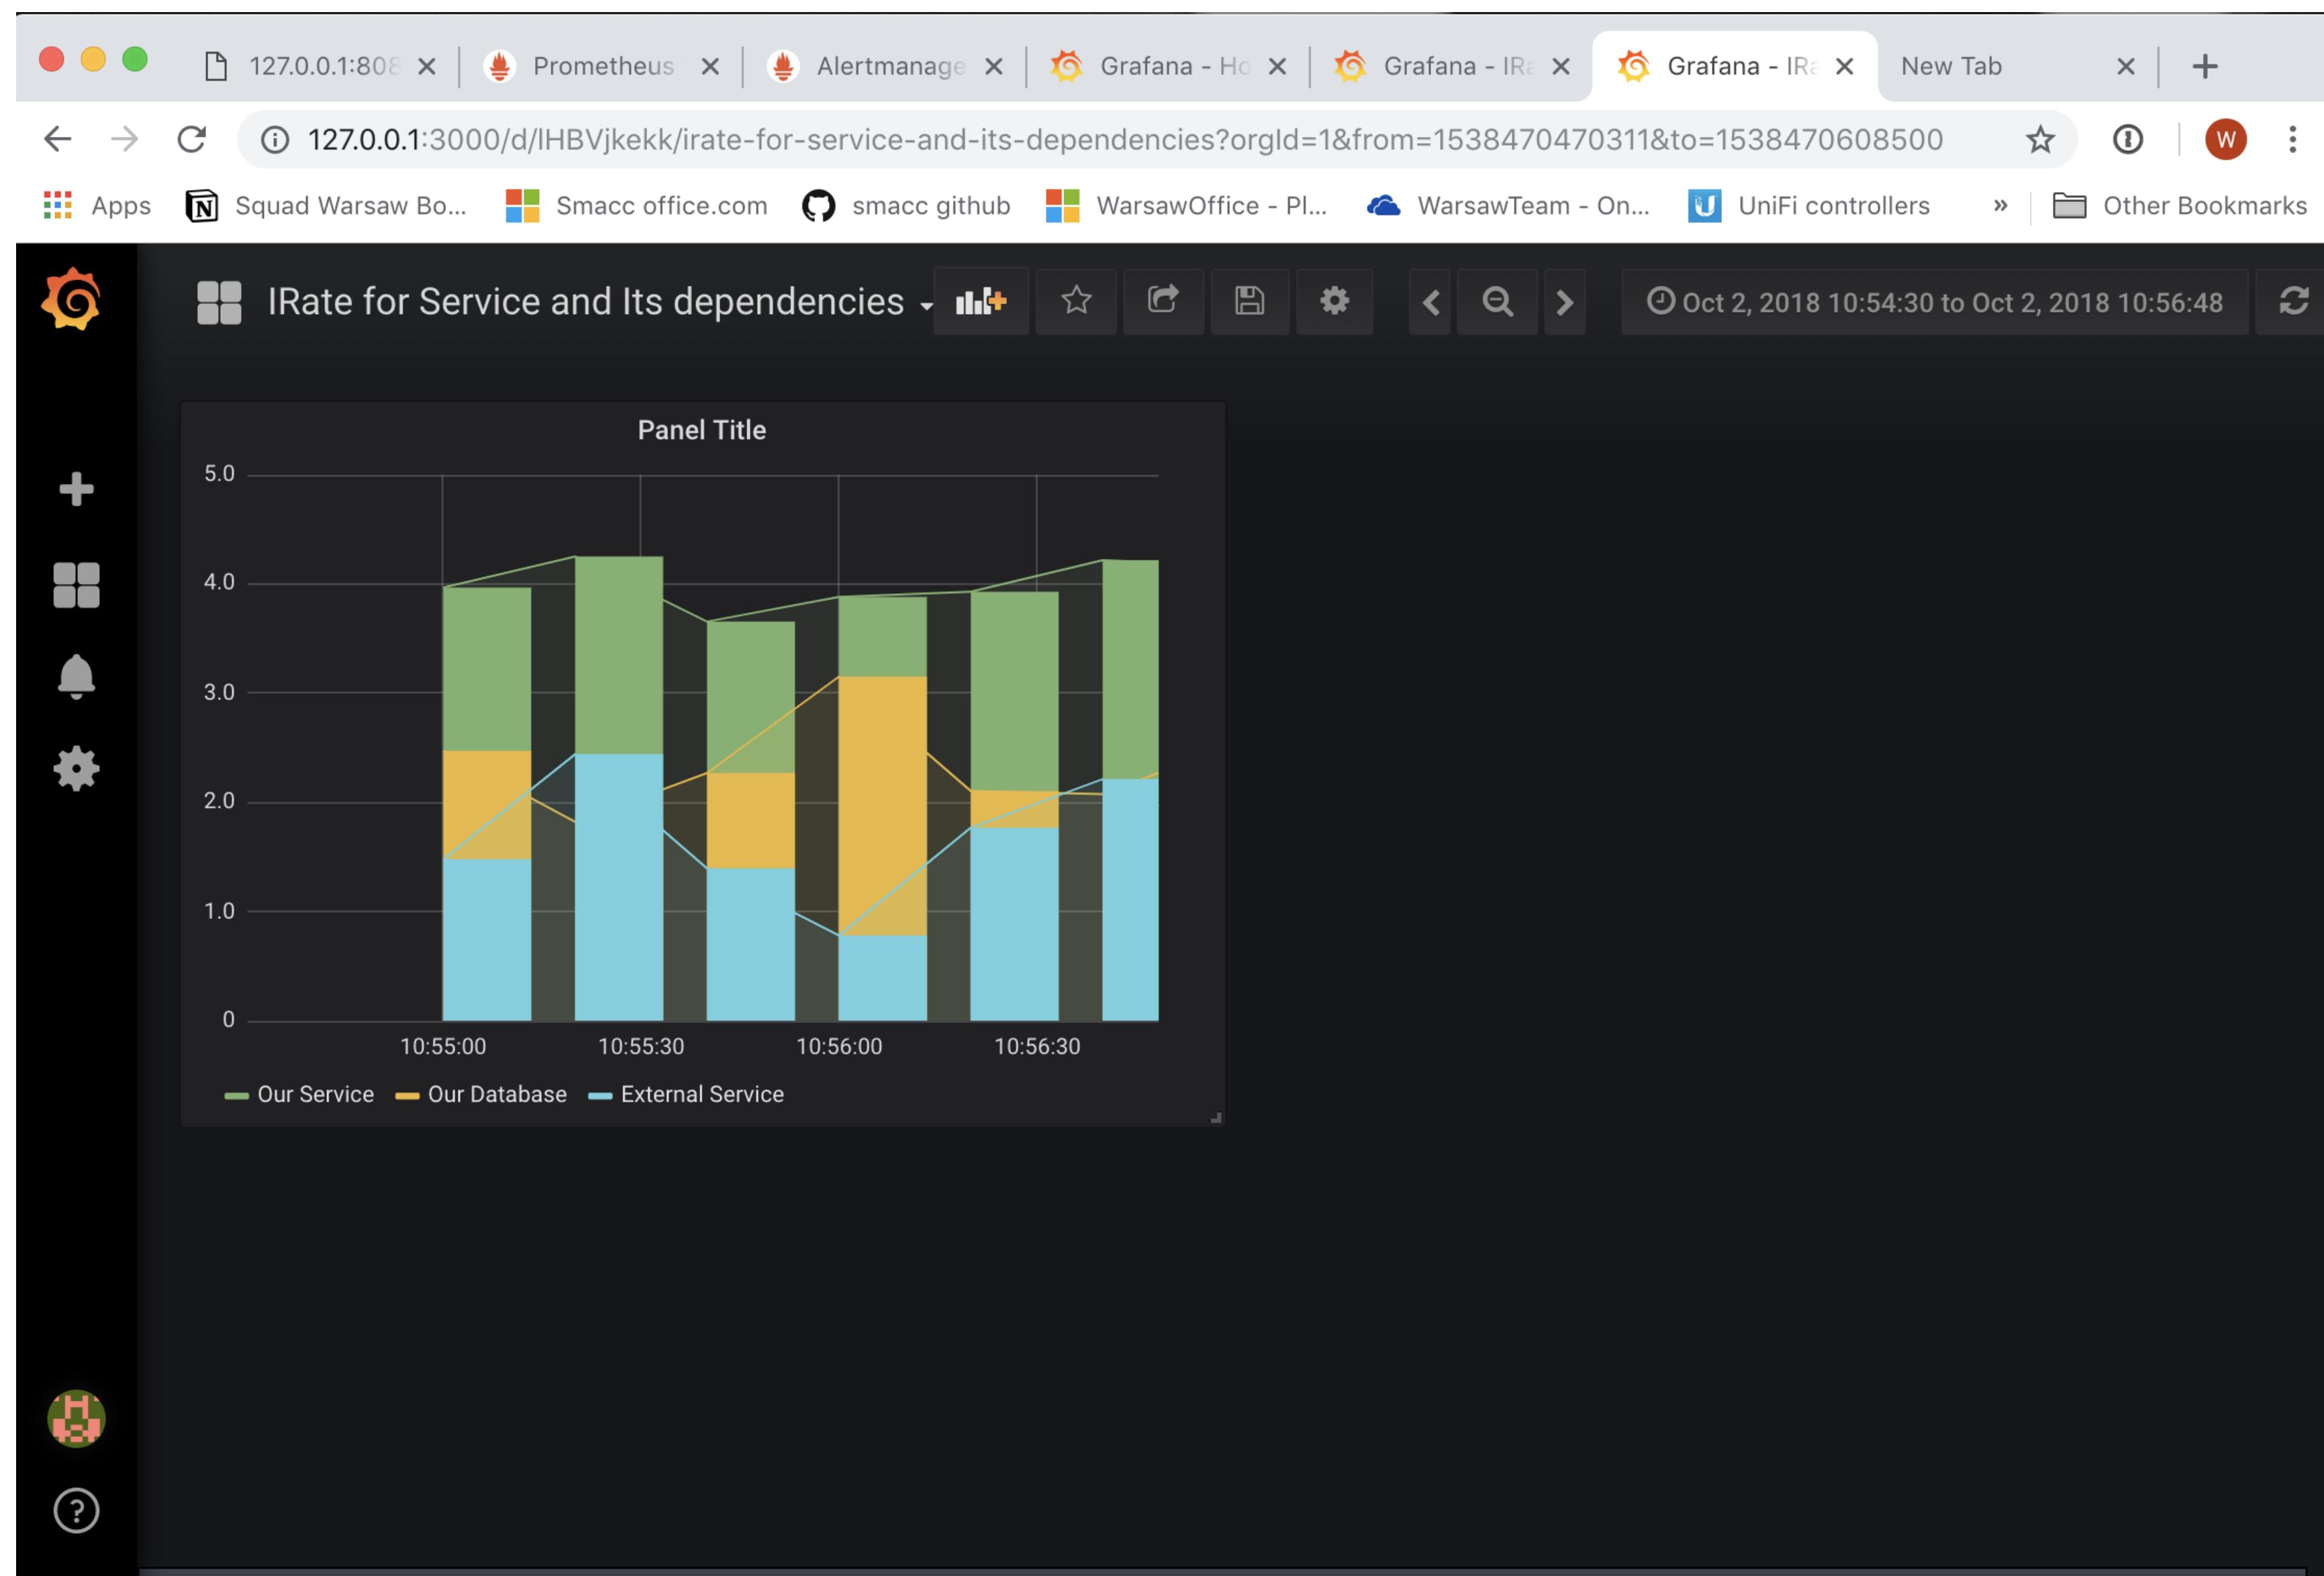
A specific alert is detailed below:

- alername="ProductionInstanceDown"** +

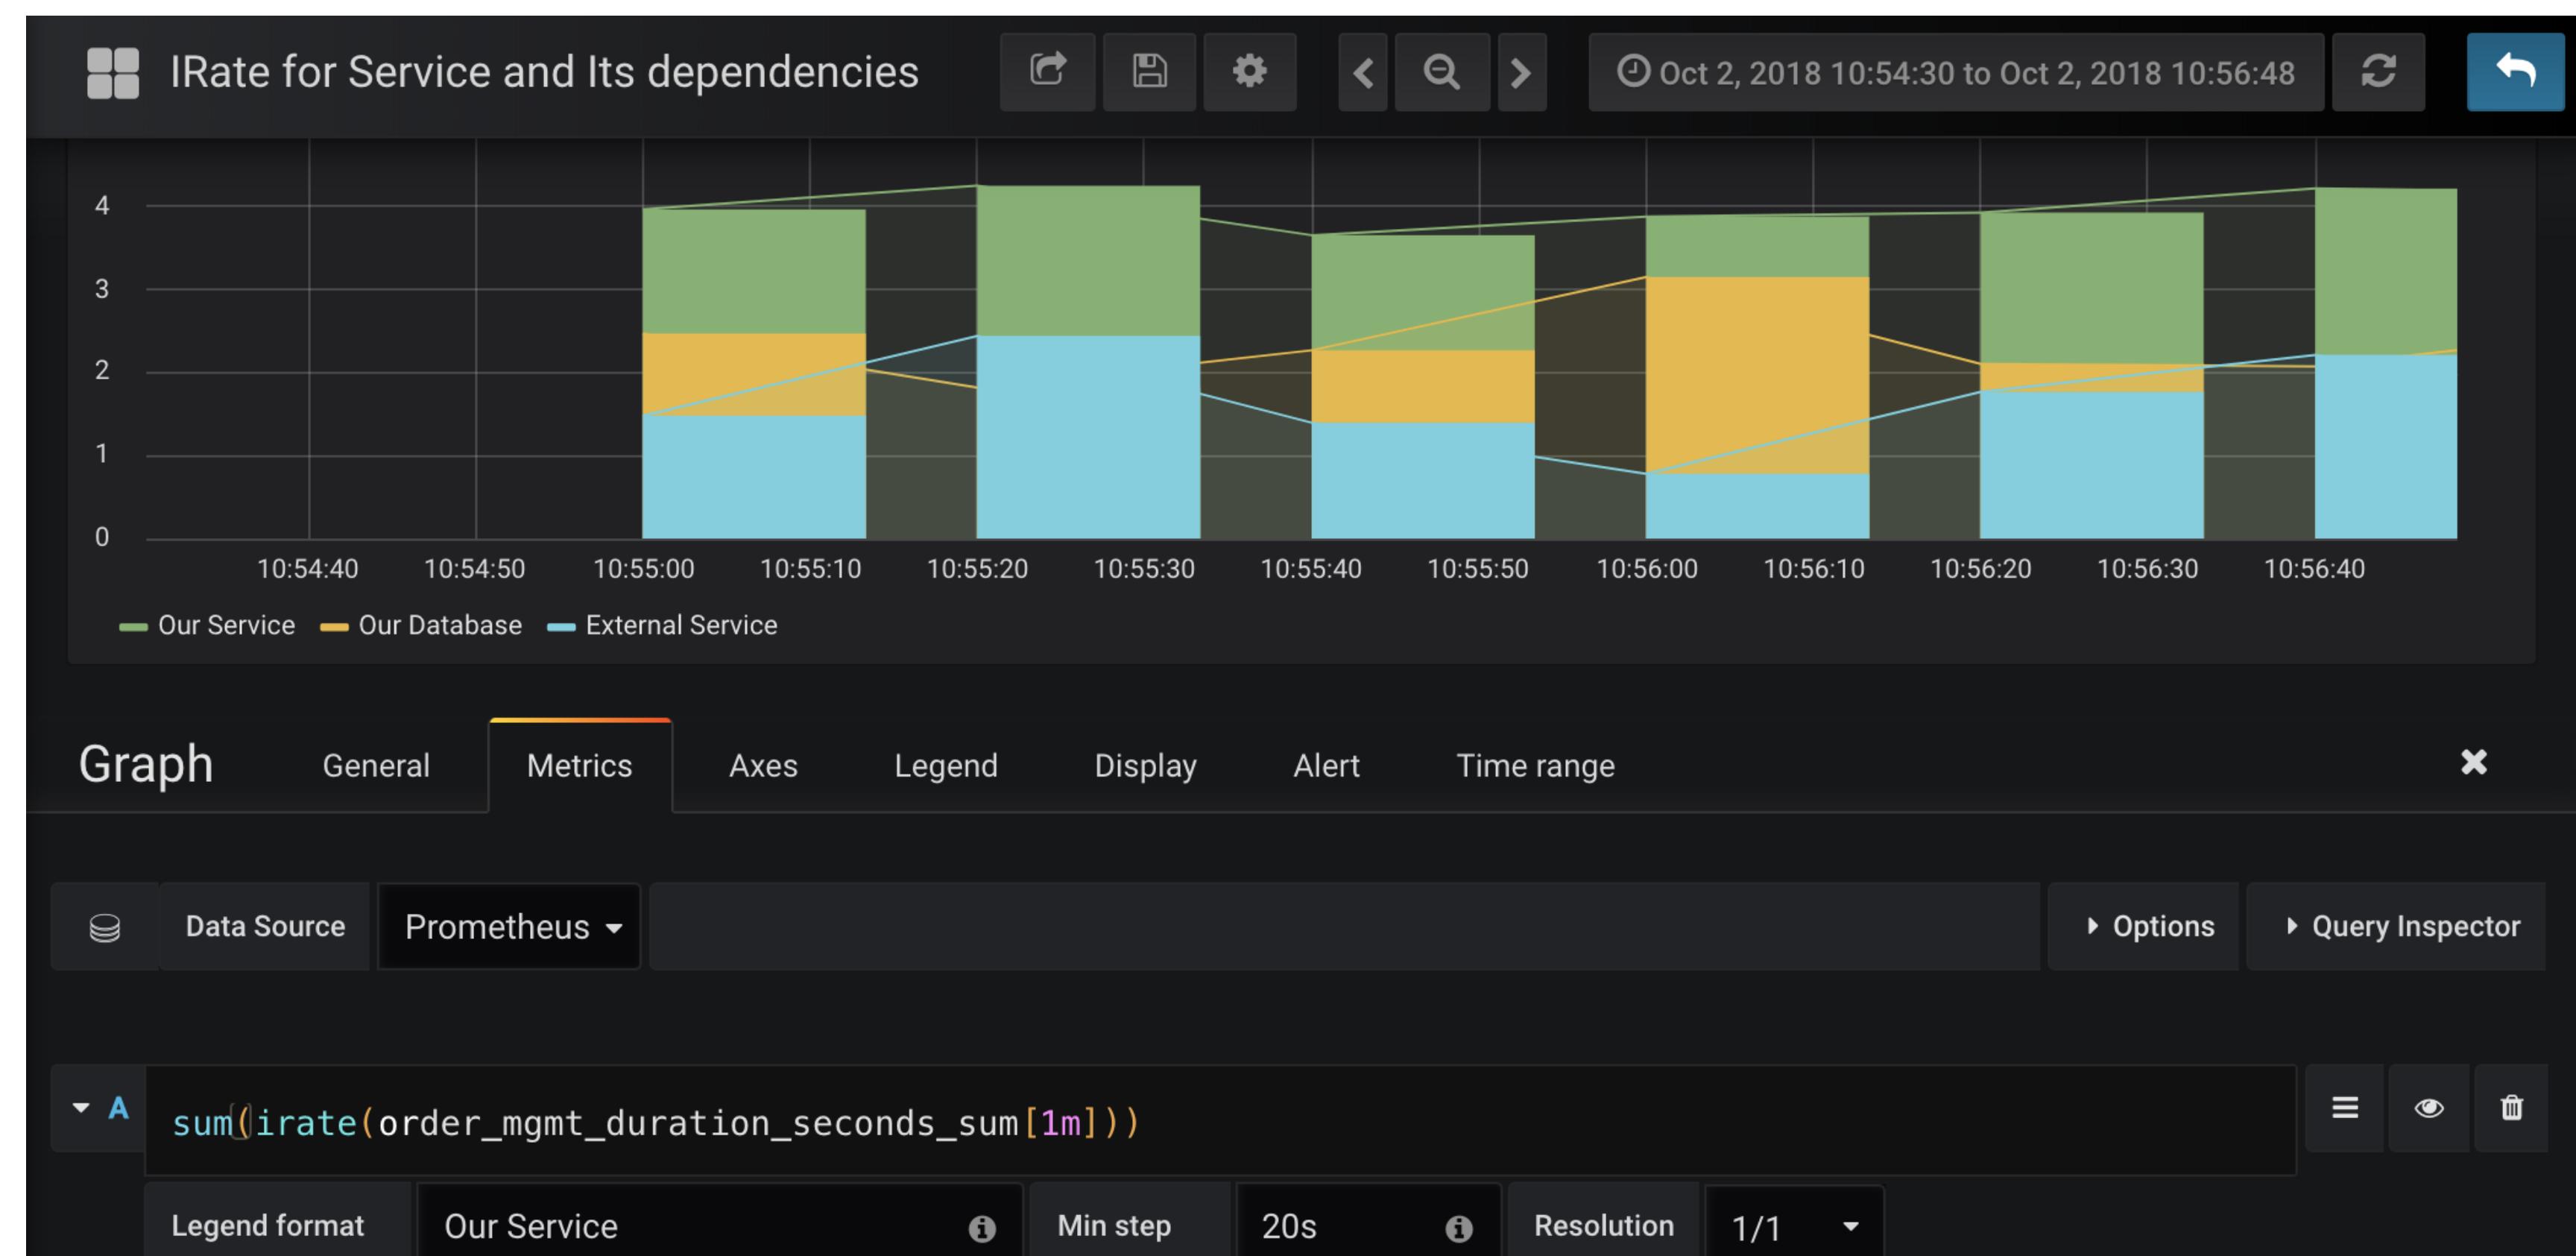
Timestamp: 08:45:01, 2018-10-02 + Info + Source + Silence

Labels:
tier="front-end" + severity="opsgenie" + owner="wb@example.com" + job="my-service" +
instance="order-manager:8080" + environment="production" + env="production" + app="order-manager" +

GRAFANA



GRAFANA



GITHUB

The screenshot shows a GitHub repository page. At the top, the URL is https://github.com/wojciech12/talk_java_2018_prometheus. The repository name is displayed in blue. A search bar and navigation links for Pull requests, Issues, Marketplace, and Explore are visible. Below the header, there's a section for managing topics. The repository stats show 13 commits, 1 branch, and 0 releases. A dropdown menu for the branch is set to 'master'. There are buttons for 'New pull request', 'Create new file', and 'Upload file'. The repository's contents are listed, including a commit by user 'wojciech12' adding a LI link in README, files for 'demo' and 'slides', a '.gitignore' file, and 'README.rst' files.

No description, website, or topics provided.

Manage topics

Branch: master ▾ New pull request Create new file Upload file

wojciech12 Add LI link in README

demo Update README.rst

slides Slides

.gitignore Demo Application

README.rst Add LI link in README

README.rst

DEMO

Everything Included:

- Prometheus dashboard and config
- AlertManager dashboard and config
- Simulate the successful and failed calls
- Examples of queries for rate

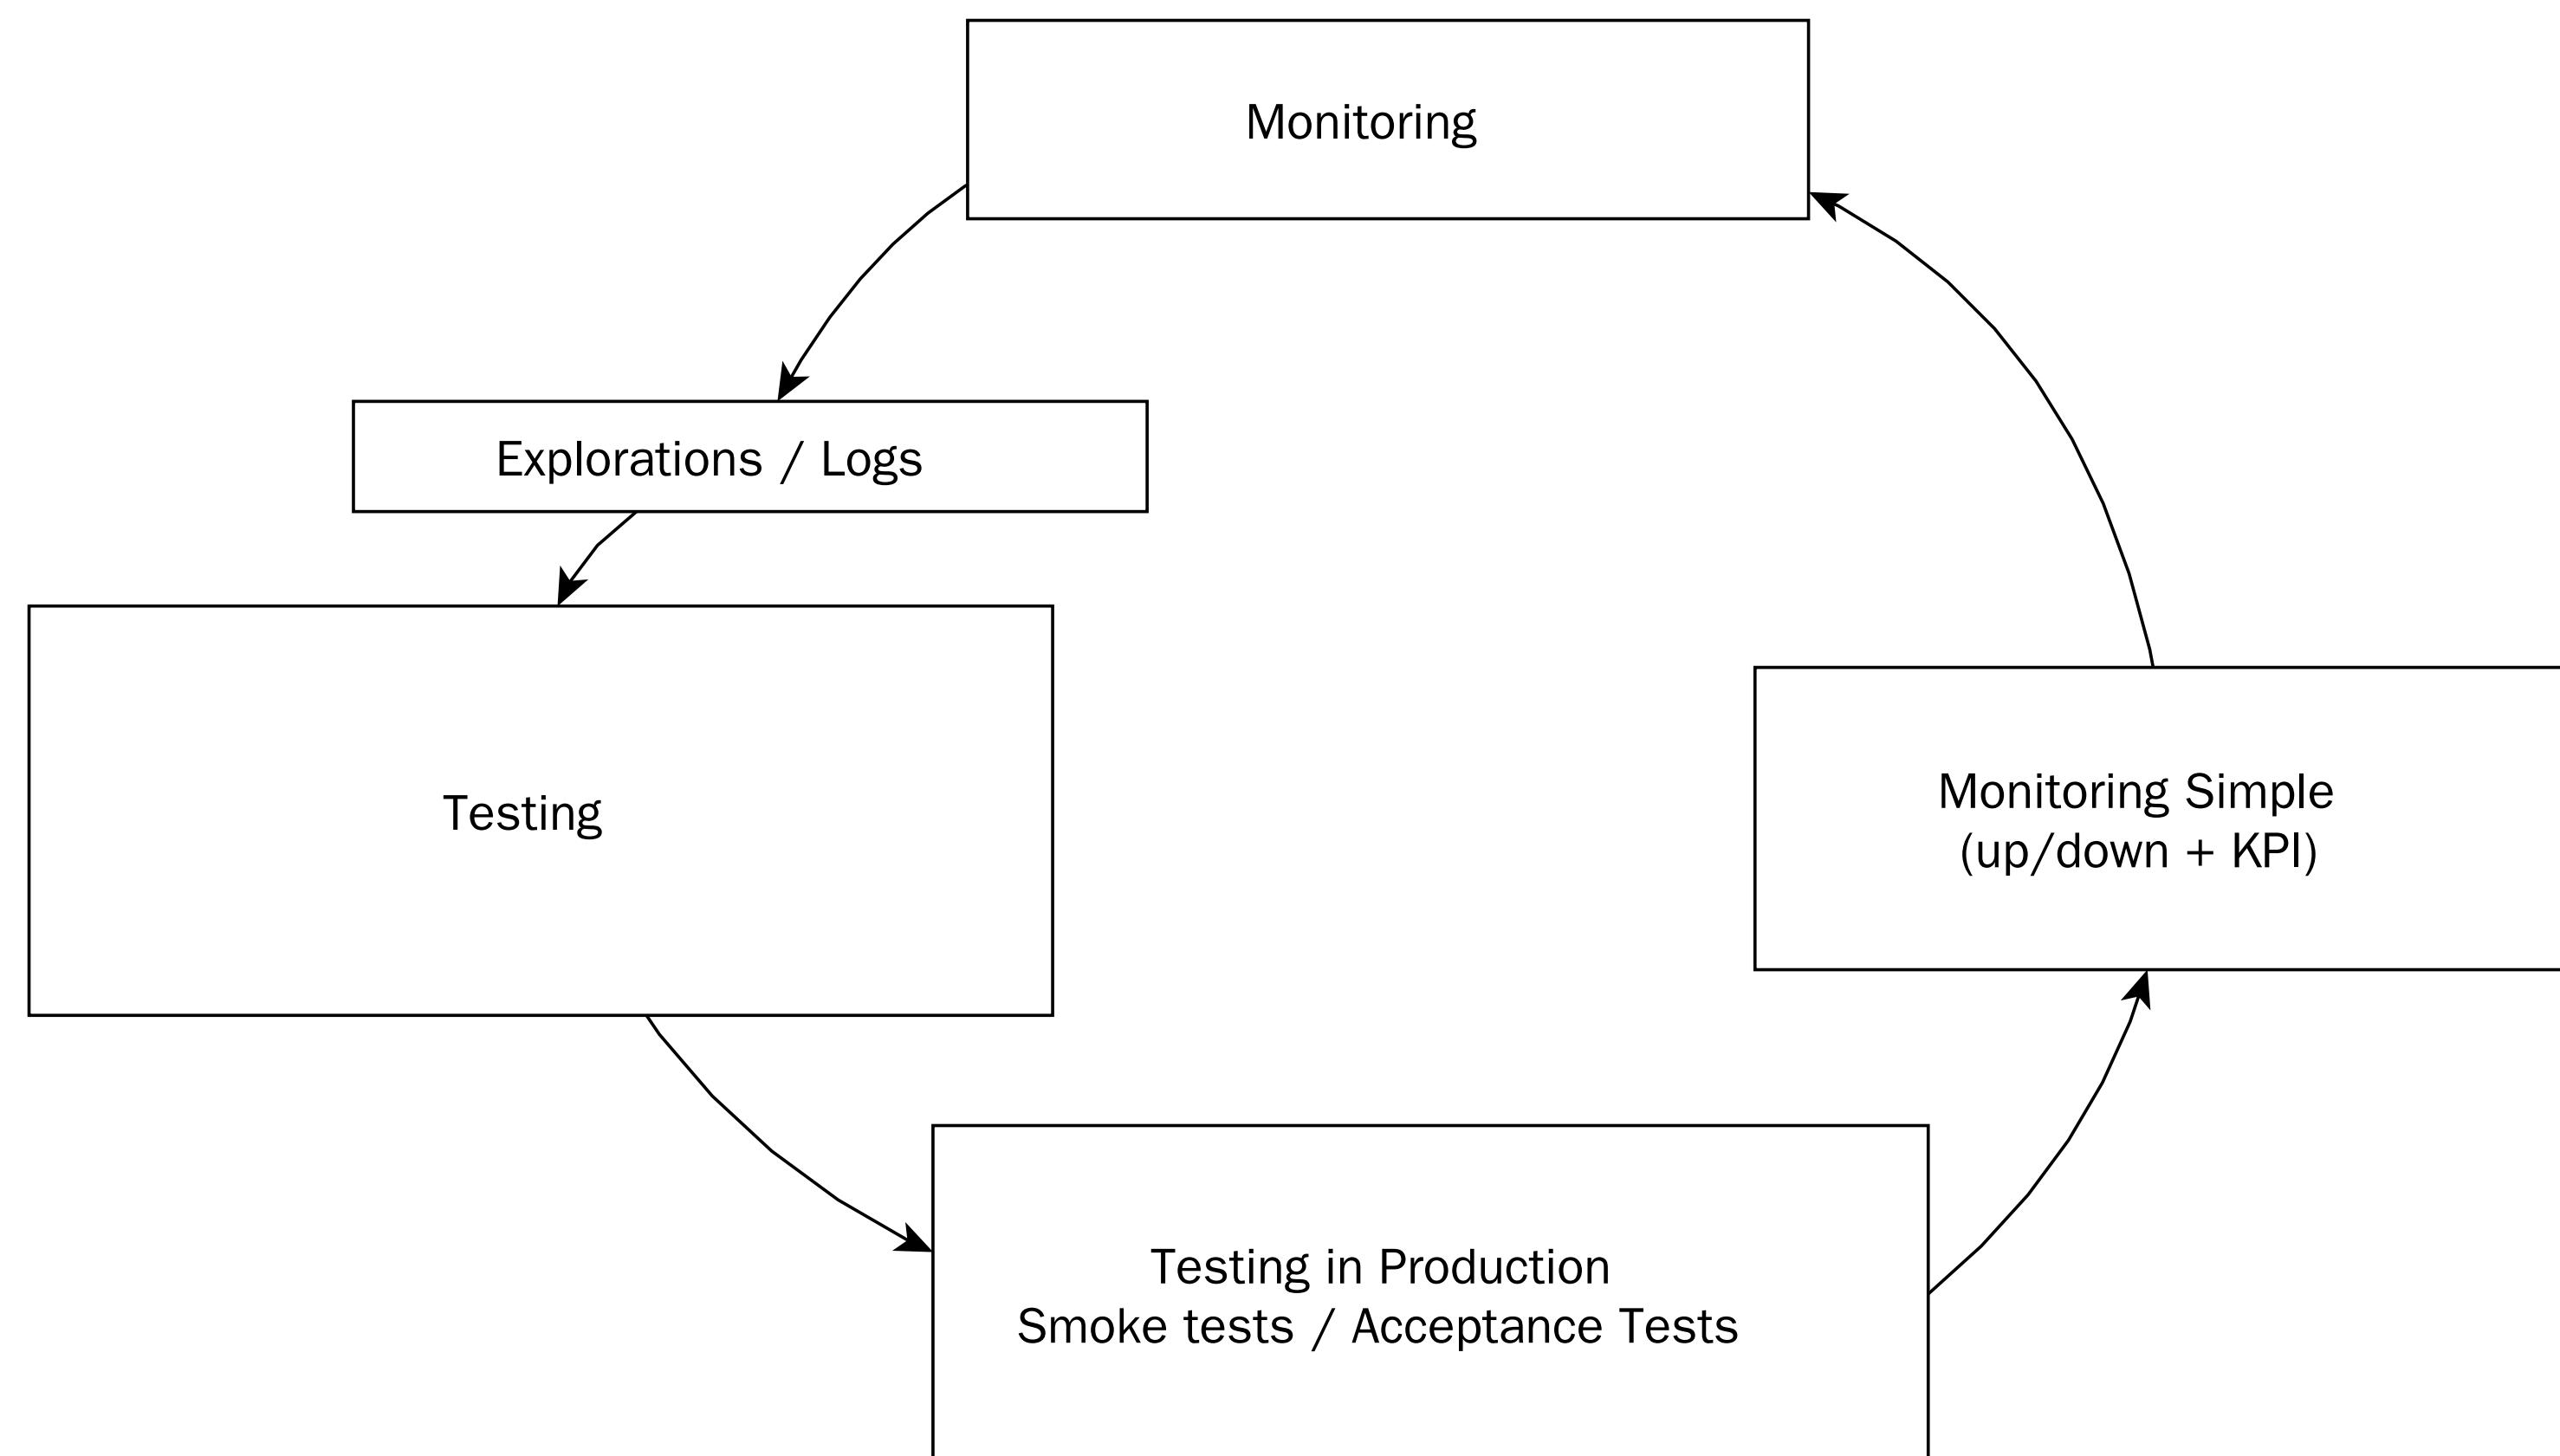
BEST PRACTISES

- Start simple (up/down), later add more complex rules
- Sum over Summaries with Q leads to incorrect results, see [prom docs](#)

SUMMARY

- Monitoring saves your time
- Checking logs == debugging vs having tests
- Logging -> high TCO

SUMMARY



THANK YOU

```
123 def distance_matrix(regions):  
124     """ Computes a distance matrix against a region list """  
125     tuples = [r.as_tuple() for r in regions]  
126     return cdist(tuples, tuples, region_distance)  
127  
128  
129 def clusterize(words, **kwargs):  
130     # TODO: write a cool docstring here  
131     db = DBSCAN(metric="precomputed", **kwargs)  
132     X = distance_matrix([Region.from_word(w) for w in words])  
133     labels = [int(l) for l in db.fit_predict(X)]
```



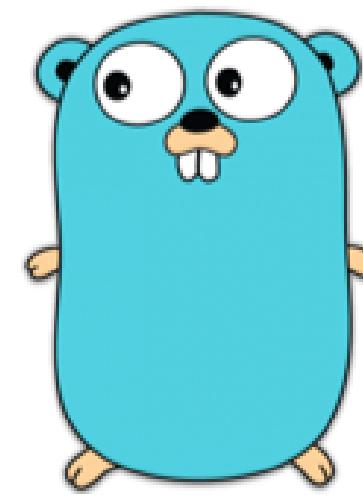
QUESTIONS?

ps. We are hiring.

```
123 def distance_matrix(regions):  
124     """ Computes a distance matrix against a region list """  
125     tuples = [r.as_tuple() for r in regions]  
126     return cdist(tuples, tuples, region_distance)  
127  
128  
129 def clusterize(words, **kwargs):  
130     # TODO: write a cool docstring here  
131     db = DBSCAN(metric="precomputed", **kwargs)  
132     X = distance_matrix([Region.from_word(w) for w in words])  
133     labels = [int(l) for l in db.fit_predict(X)]
```



SMACC



Go



PYTORCH

TensorFlow™



amazon
web services™

Azure

λ



BACKUP SLIDES

```
123 def distance_matrix(regions):  
124     """ Computes a distance matrix against a region list """  
125     tuples = [r.as_tuple() for r in regions]  
126     return cdist(tuples, tuples, region_distance)  
127  
128  
129 def clusterize(words, **kwargs):  
130     # TODO: write a cool docstring here  
131     db = DBSCAN(metric="precomputed", **kwargs)  
132     X = distance_matrix([Region.from_word(w) for w in words])  
133     labels = [int(l) for l in db.fit_predict(X)]
```



PROMETHEUS - PUSH MODEL

- See:

<https://prometheus.io/docs/instrumenting/pushing/>

Good for short living jobs in your cluster.

DESIGNING METRIC NAMES

Which one is better?

- `request_duration{app=my_app}`
- `my_app_request_duration`

see documentation on best practises for [metric naming](#) and [instrumentation](#)

DESIGNING METRIC NAMES

Which one is better?

- `order_mgmt_db_duration_seconds_sum`
- `order_mgmt_duration_seconds_sum{dep_name='db'}`

PROMETHEUS + K8S = <3

**LABELS ARE PROPAGATED FROM K8S TO
PROMETHEUS**

INTEGRATION WITH PROMETHEUS

```
cat memcached-0-service.yaml
```

```
---  
apiVersion: v1  
kind: Service  
metadata:  
  name: memcached-0  
  labels:  
    app: memcached  
    kubernetes.io/name: "memcached"  
    role: shard-0  
  annotations:  
    prometheus.io/scrape: "true"  
    prometheus.io/scheme: "http"  
    prometheus.io/path: "metrics"  
    prometheus.io/port: "9150"
```

<https://github.com/skarab7/kubernetes-memcached>